HANDOUT B

Why does this matter and how do you use it?

Today, answering "is this model good enough for your workload?" means running evaluation on hundreds of items per model, per cycle. The *Spectral-IRT Stratified Pruner* cuts that to 72 items for coding, 24 for long-context reasoning, and 42 for multimodal, returning the same go/no-go answer in a fraction of the time. For a sales engineer, this means a customer can get a verified, defensible answer the same day rather than waiting for a full evaluation run.

To run it tomorrow, point *evalscope* at the stored predictions and reviews for the candidate model and invoke the pruner. The output is a go/no-go flag and a score for each model, and no statistics background is required to act on it. The pruned index is reusable, so the same subset can be applied to every new model without rebuilding it from scratch.

The multimodal probe does something random sampling cannot. Charts, medical images, diagrams, and photographs each require different visual capabilities, and a model can score well on average while failing silently on one type. Those are precisely the failure modes random sampling has no mechanism to surface. This algorithm is designed to expose them, giving the customer a targeted answer about visual reasoning quality rather than a general accuracy number.

For a customer-facing PM, both pruners compress the evaluation cycle without compressing confidence in the answer. Faster evaluations mean shorter sales cycles, quicker iteration when a customer's workload changes, and a concrete artifact, the pruned index, that can be handed to a deployment lead as the agreed testing standard for that customer's use case.